




ChatGPT

*Helping users evaluate AI outputs
before they trust them.*

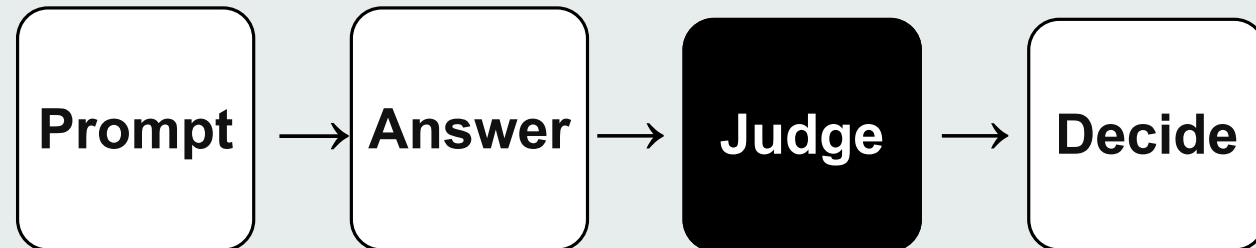


Prompt → Answer → Judgment
not blind trust

AI output evaluation is now a trust problem in high-stakes ChatGPT workflows

What is AI output evaluation?

Users must judge quality, not just read.



The trust gap

What users see ?

Polished, confident, complete looking answer

What actually there?

Missing context, weak reasoning, hidden assumption

The result

Users have no way to detect the gap

what is unmet needs ?

- Users lack visibility into:
- Missing context.
- Hidden assumptions
- Uncertainty
- weak reasoning
- Alternative viewpoints

Market signals

Adoption is scaling faster than users' evaluation habits.

700M+

Weekly ChatGPT users by July 2025

63%

Surveyed users use ChatGPT daily,

88%

Organizations regularly use AI

Key drivers for growth?

- AI becoming embedded in work and education workflows
- AI becomes a decision-support system,
- Users are moving from using AI for information retrieval to brainstorming, analysis etc

What remains unsolved?


- Users still judge by fluency and confidence.
- “Good enough” depends on the task.
- Users rarely build evaluation skill over time.

Opportunity Area

Help users evaluate before they trust, while keeping judgment in control.

Research Insights

How current model works?



What it leads to?

Weak output used → rework / error → unstable trust

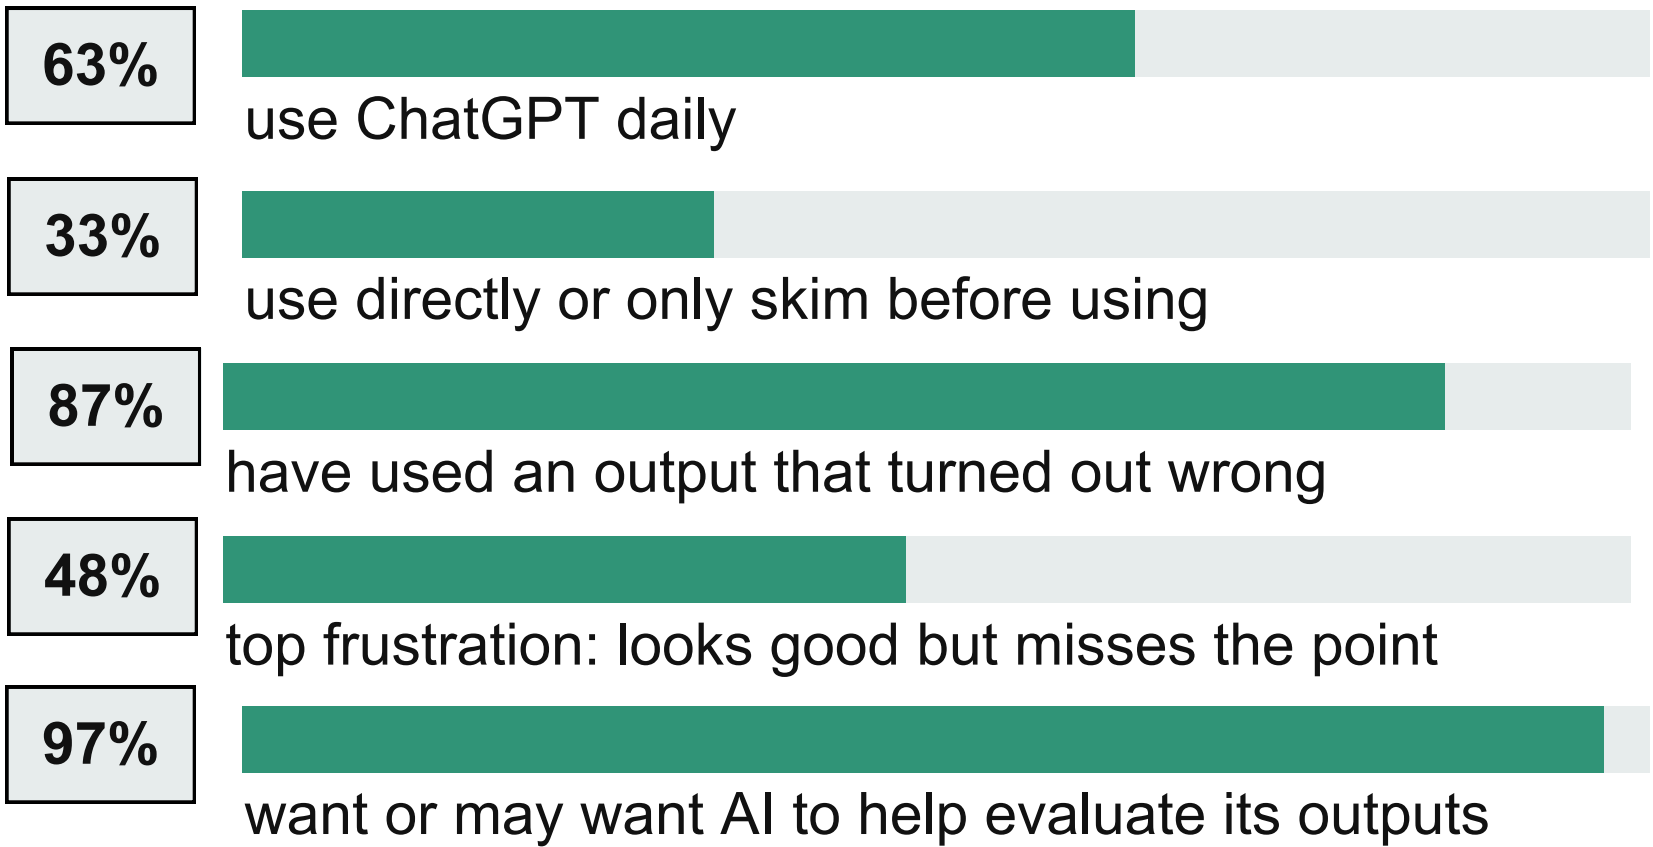
Hypotheses: major problems users are facing

- Users trust polished outputs without evaluating because they look complete and confident.
- Users don't realize ChatGPT made assumptions they never see or were asked about.
- Missing context and assumptions stay invisible until the answer fails in real use.
- ChatGPT sounds equally confident about everything even at highly uncertain things.

◆ Primary research

Quantitative research

Key signals from survey responses



User interview n=4 Survey n=30 Survey response [🔗](#)

Qualitative signals

What users asked ChatGPT to improve

- 1 Clearer uncertainty**
Users asked for signals like “I am not fully sure” and “you should verify this.”
- 2 Proof before trust**
Responses should validate claims with sources, references, or evidence when possible.
- 3 Reasoning before action**
Users want ChatGPT to think through edge cases instead of jumping directly to output.
- 4 Intent matching**
Trust drops when the answer sounds good but does not match the user’s actual prompt.

Key Takeaways Users want help evaluating answers, but they do not want ChatGPT to become the final authority. The product should expose gaps and keep judgment with the user.

Secondary research: competitor analysis

Product	Current approach	Strength	User pain-points
ChatGPT	Search, reasoning and broad workflows	Widest user base , plugin ecosystem	still missing context and quality , No uncertainty shown
Claude	Long-context writing, coding and artifacts	Good for coding , designin tools integration	Difficult to assess assumptions, completeness
Perplexity	Real-time answer engine with cited sources	Fast source-checking for factual claims , Source transparency	Don't help users judge reasoning quality or completeness
Gemini	Google Search plus Workspace integrations	Fact verification , google workspace	Fact-checking alone doesn't build judgment , Limited support for understanding uncertainty

User Persona



Riya, 23

Final-year college student
Uses ChatGPT for research, summaries, resumes and career decisions.

“It sounds good, but I do not know if I should trust it.”

Pain-points

- Polished answers hide weak reasoning or missing context
- Manual verification slows down every important task

Goals

- Know what is safe to use, refine or verify
- Need to Get simple explanations and guidance when trying to understand something for the first time.

User segment

Students & Early Professionals
Age 18-27 | Tier 1 & 2 cities | India

- Use cases** - learning, research, writing, resume/career prep ,assignment & professional tasks
- Behaviour**- Rely heavily on AI for research they can't self-verify. Judge output by how polished it looks not how accurate it is. and use outputs directly.

Why this segment?

- High stakes tasks - wrong research = failed assignments and bad decisions
- Low domain expertise - can't spot errors themselves unlike doctors or lawyers
- No evaluation framework - they have nothing to check against
- Heaviest ChatGPT users - 63% use it daily (from our survey)
- Growth opportunity - convert free users to paid by building trust

User Persona



Mba student / aspiring investment banker

Uses ChatGPT for Assignment ,
resumes and Preparing for interviews
and placements

Abhinav, 24

Pain-points

- Doesn't know if outputs are good enough
- Conflicting information across sources

Goals

- Learn faster and fill knowledge gaps quickly
- Create impactful presentations, reports, and case study submissions

Problem framing canvas

What is the true problem?

Users are not failing only because ChatGPT can be wrong. They struggle to judge whether a polished answer is correct, complete, useful and decision-ready.

Core problem: perceived quality often feels higher than actual quality.

Value for target customers

- Builds trust in ChatGPT - Users trust more because they understand more
- Fewer weak outputs used in real work
- Fills the evaluation gap - Users can see what's uncertain, what's missing, and what to verify before acting

Who is facing this problem?

Students and early-career professionals using ChatGPT for learning, research, writing, resume and for professional work

High AI adopters who care about getting tasks done efficiently but lack the domain expertise to verify outputs.

Value for the business

- Increased retention - Users who trust outputs stay longer and churn less
- Higher paid conversion - Building trust converts free users to Plus subscribers
- Competitive moat - No competitor helps users evaluate this differentiates ChatGPT

How do we know it is real?

- **86.7%** used an AI output that was wrong
- **63.3%** verify high-stakes tasks often
- **46.7%** say outputs can look good but miss the point

Why solve this problem now?

- AI adoption is accelerating - More users doing high-stakes work on ChatGPT every day
- Trust gap is widening - As outputs get more polished, the gap between perceived and actual quality grows
- No one is solving this - No competitor addresses output evaluation — first mover advantage
- Users are asking for it - 97% of our survey said they want help evaluating outputs

Solution Ideation & Prioritization

After ChatGPT gives an answer, users can choose a review mode when they are unsure, dissatisfied, or using the output for an important task.

User decides

Prompt → Answer → Review Mode → Trust / Refine / Verify / Use

Solution 1: Review Mode

Fast trust calibration after the answer

- Shows assumptions made by ChatGPT
- Surfaces missing information/context
- Lists areas users should verify

Solution 2: Goal-Based Checklist

Quality depends on the user’s goal

- User selects goal: learning, research, coding, career
- Shows what a good answer should include
- Marks what is present vs missing

Solution 3: Blind Spot Finder

Find what the answer may have ignored

- Identifies ignored risks and missing perspectives
- Surfaces unasked questions and weak assumptions
- Reduces one-sided or generic outputs

[Solution link](#)


Prioritization using ICE scoring

Solution	Impact	Confidence	Ease	Why it matters	ICE
Review Mode	9	9	8	Directly makes assumptions, missing context and verification visible	648
Goal-Based Checklist	8	8	7	Teaches users what quality means for their specific goal	448
Blind Spot Finder	8	7	7	Reveals ignored risks and one-sided thinking before use	392

Chosen direction: Review Mode it is the clearest reflection layer between a polished answer and the user’s decision.

Review Mode adds a reflection layer between ChatGPT’s answer and the user’s decision

User data flow



User journey mapping


Step	User action	User emotion	Pain-point solved
Intent	Asks a complex / important question	Urgent: “I need a good answer fast”	Starts with normal ChatGPT speed
Receive	Gets a polished answer	Unsure: “This looks right, but is it?”	Separates polish from trust
Review	Clicks “Review this answer”	Curious: “What did it assume?”	Makes assumptions visible
Calibrate	Reads missing info + verify areas	Clear: “I know what to check”	Surfaces gaps and uncertainty
Act	Refines, verifies, or uses answer	In-control: “I decide what to trust”	Supports human judgment

Key behavior shift

From “answer looks polished, so I’ll use it” → “I understand what is assumed, missing, and worth verifying before I act.”

Prototype: answer → review → refine / verify / use

Wireframing flow



Links :

[Prototype](#)


[wireframe](#)

[Test prompt](#)

[System Architecture](#)

Metrics

Metric logic



North Star Metric
% of high-context sessions where users take a calibrated next step
Refine prompt • verify claim • use with confidence

Guardrail metrics
Review fatigue / dismiss rate
Drop-off after answer generation
Time-to-useful-output remains acceptable
Users do not treat review output as final authority
No rise in excessive skepticism or reduced productivity

Leading Indicators

Category	Metric	What it measures
L0	Review trigger CTR	Whether users notice and intentionally open evaluation
L1	Review completion rate	Whether users inspect assumptions, missing info and areas to verify
L1	Prompt refinement after review	Whether users add missing context instead of blindly regenerating
L1	Verification action rate	Whether users follow through on facts or claims that need checking

Lagging Indicators

Category	Metric	What it measures
L0	Calibrated decision rate	% of reviewed sessions ending in trust, refine, verify or use with clarity
L1	Repeat review usage	Whether the toolkit becomes a trusted evaluation habit
L1	Reduced low-quality regenerations	Fewer repeated prompts caused by uncertainty or unclear gaps
L2	High-stakes workflow retention	More users return for research, career, writing and analysis tasks

◆ Monetization plan

Business model

Free

7 reviews/week for high-context tasks; upsell when users ask for deeper review.

Plus / Pro

Higher review limits, Toolkit, saved review history and reusable review templates.

Team

Policy-based review triggers, workspace templates and admin-safe evaluation patterns.

Key trade-offs

Speed

vs

Reflection

Caution

vs

Productivity

Assistance

vs

Dependence

◆ Risks and mitigation

1

Risk: Review becomes another black box

Mitigation: No trust scores; show concrete assumptions, missing info and areas to verify.

2

Risk: Users feel slowed down

Mitigation: Trigger only on high-context tasks and keep a clear skip path.

3

Risk: Users blindly trust the review

Mitigation: Frame as guidance, not verdict; user chooses trust, refine, verify or use.

4

Risk: Too much transparency creates overload

Mitigation: Start with 3 short sections; deeper detail only on click.

5

Risk: Users become overly skeptical

Mitigation: Label outputs as useful starting points; do not make every answer feel risky.

Distribution model

- **In-product trigger:** “Review before using” under high-context answers
- **Onboarding nudges:** students, research, writing and career prep journeys
- **Template gallery:** resume review, paper summary, decision advice, content drafts
- **Growth loop:** shareable reviewed outputs with assumptions + verification preview

Business impact

Trust → Retention → Deeper usage

- Safer adoption for research, writing, career prep and analysis
- More refinement loops: answer → review → improve → use
- Differentiates ChatGPT as a thinking partner, not just an answer engine